Million Dollar Baby An 'angr'y attempt at the Cyber Grand Challenge

Nick Stephens

stephens@cs.ucsb.edu



Shellphish



- Who are we?
 - A team of security enthusiasts
 - do research in System Security
 - play Capture the Flag competitions
 - released a couple of tools

Mini-primer: What's a CTF?



- Security competition:
 - exploit a vulnerable service / website / cryptosystem
 - reverse a binary
 - -
- Different formats
 - Jeopardy Attack-Defense
 - Online Live
 - -
- Basic idea: find the secret, submit to organizers, ... profit

Shellphish





Roadmap



- DARPA Cyber Grand Challenge (CGC)
 The (almost-)Million Dollar Baby
- Our Cyber Reasoning System (CRS)
 Fancy term for auto-playing a CTF
- Automated Vulnerability Discovery
 Driller
- Automated Vulnerability Exploitation
 How it works
- Auto-exploitation demo using angr
 Open-source binary analysis framework

Cyber Grand Challenge (CGC)



- 2004: DARPA Grand Challenge
 - Autonomous vehicles



Cyber Grand Challenge (CGC)



- 20**14**: DARPA **Cyber** Grand Challenge
 - Autonomous hacking!



Cyber Grand Challenge (CGC)



- Qualification event: ONLINE
 - ~70 teams → 7 qualified teams
 - \$750k per team



- Final event: DEFCON 24
 - Winning CRS will also play against humans!

cybergrandchallenge.com / cgc.darpa.mil



CGC Rules



- Attack-Defense CTF
- Solving security challenges → Developing a system that solves security challenges
- Develop a system that automatically
 - Exploits vulnerabilities
 - Removes vulnerabilities
- No human intervention

CGC – Rules



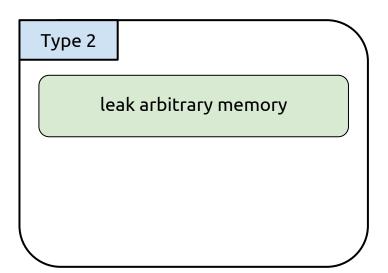
- Exploits
 - "getting a flag" (how? where?)
 - For the quals: exploit = **crash**
 - For the finals: exploit = **demonstrate PC control**
- Defend
 - int main() { return 0; }
 - Functionality checks
 - SIGSEGV => exit(0)
 - No easy "out-of-band" error handling
 - QEMU-style interpreter: interrupts => exit(0)
 - Performance cost (CPU, memory, file size)

CGC Final Round Exploits



control arbitrary register

control program counter



CGC Qualification Event – Rules



- Basic idea:
 - Real(istic) programs
 - No "extra" complications

CGC Qualification Event – Rules



- Architecture: Intel x86, 32-bit
- OS: DECREE
 - Linux-like, but with 7 syscalls only
 - transmit / receive / fdwait (≈select)
 - allocate / deallocate (even executable!)
 - random
 - _terminate
- no signals, threads, shared memory
- "Bring Your Own Defense" approach (and pay for it)
 - Not even "the usual": stack is executable, no ASLR, ...

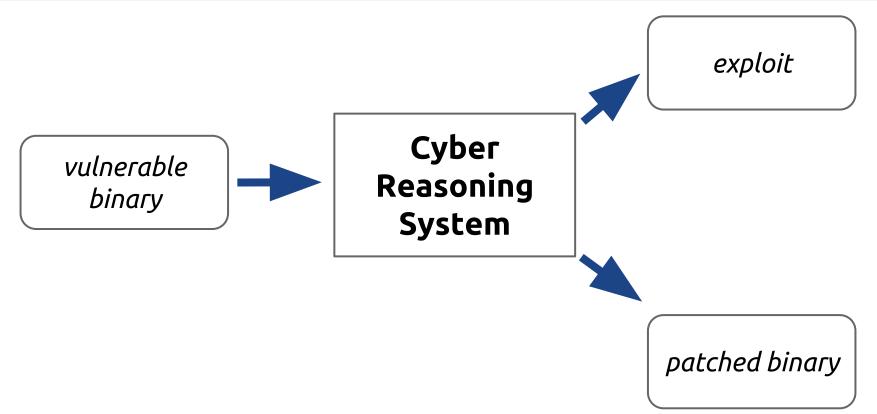
Roadmap



- DARPA Cyber Grand Challenge (CGC)
 The (almost-)Million Dollar Baby
- Our Cyber Reasoning System (CRS)
 Fancy term for auto-playing a CTF
- Automated Vulnerability Discovery
 Driller
- Automated Vulnerability Exploitation
 How it works
- Auto-exploitation demo using angr
 Open-source binary analysis framework

Shellphish CRS

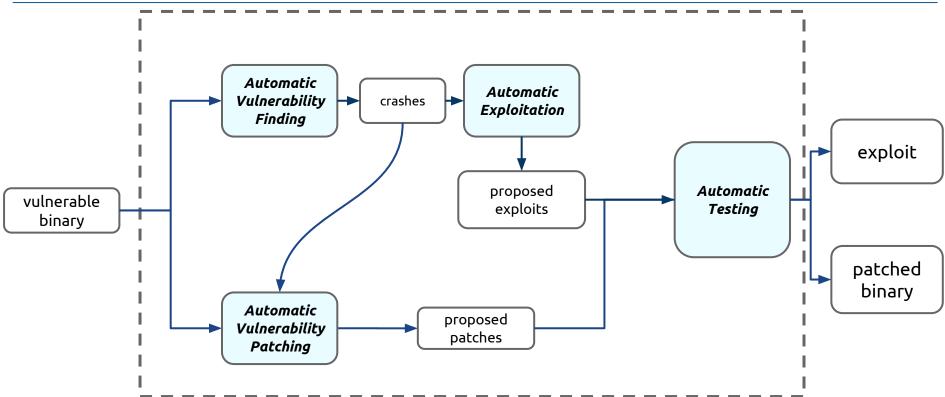




16

Shellphish CRS





17

Roadmap



- DARPA Cyber Grand Challenge (CGC)
 The (almost-)Million Dollar Baby
- Our Cyber Reasoning System (CRS)
 Fancy term for auto-playing a CTF
- Automated Vulnerability Discovery
- Automated Vulnerability Exploitation
 How it works
- Auto-exploitation demo using angr
 Open-source binary analysis framework

Quals Approach to Finding Bugs



Fuzz a binary

... and ...

Symbolically explore a binary

Finding bugs



Fuzzing

```
x = int(input())
if x > 10:
    if x < 100:
        print "You win!"
    else:
        print "You lose!"
else:
    print "You lose!"</pre>
```





Catching Bugs



- Monitors program for crashes

```
x = int(input())
if x > 10:
    if x^2 == 152399025:
        print "You win!"
    else:
        print "You lose!"
else:
    print "You lose!"
```

Let's fuzz it!



```
1 ⇒ "You lose!"
593 ⇒ "You lose!"
183 ⇒ "You lose!"
4 \Rightarrow "You lose!"
498 ⇒ "You lose!"
42 ⇒ "You lose!"
3 \Rightarrow "You lose!"
```

 $57 \Rightarrow$ "You lose!"

Finding bugs

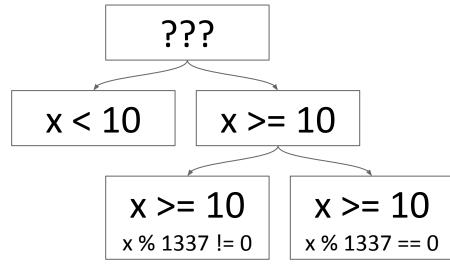


Symbolic Execution

```
x = input()
  if x > = 10:

    if x % 1337 == 0:

                              x < 10
        print "You win!"
   clse:
        print "You lose!"
clse:
     print "You lose!"
```



```
555
x = input()
if x >= 10:
   if x % 1337 == 0:
                               x < 10
                                          x >= 10
      print "You win!"
   else:
                                                x >= 10
                                    x >= 10
      print "You lose!"
                                    x % 1337 != 0
                                                x % 1337 == 0
else:
   print "You lose!"
```

Catching Bugs



- Checks each state for safety violations
 - symbolic program counter
 - writes/reads from symbolic address

```
x = input()
def recurse(x, depth):
  if depth == 2000:
    return 0
 else:
    r = 0;
   if x[depth] == "B":
      r = 1
    return r + recurse(x[depth], depth)
                                                       ???
if recurse(x, 0) == 1:
 print "You win!"
else:
                                          x[d] != "B"
                                                            x[d] == "B"
 print "You lose!"
```

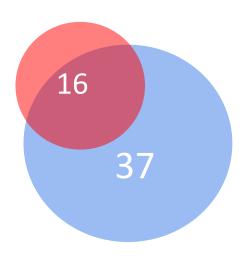
Quals Performance



Symbolic Execution (angr): 16

Fuzzing (AFL): 37

Total: 44 / 128



... need something better

Driller



Driller: Augmenting Fuzzing through Selective Symbolic Execution

NDSS 2016

Different Approaches



Fuzzing

 Good at finding solutions for general conditions

 Bad at finding solutions for specific conditions

Symbolic Execution

 Good at finding solutions for specific conditions

 Spends too much time iterating over general conditions

Fuzzing vs. Symbolic Execution



```
x = input()
def recurse(x, depth):
  if depth == 2000:
    return 0
  else:
    r = 0:
    if x[depth] == "B":
      r = 1
    return r + recurse(x[depth], depth)
if recurse(x, 0) == 1:
  print "You win!"
```

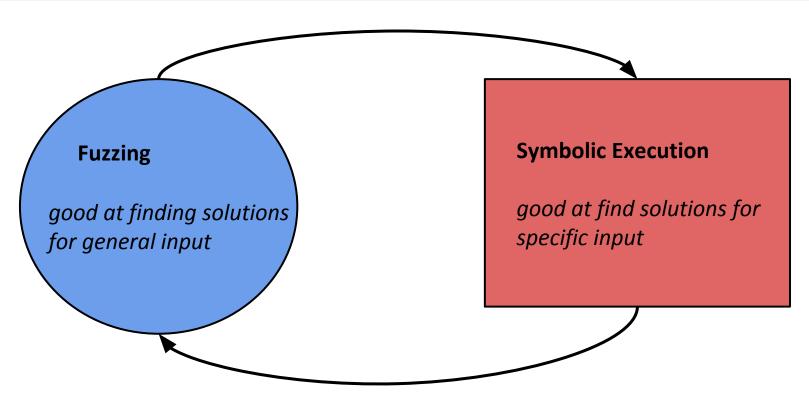
```
x = int(input())
if x >= 10:
    if x^2 == 152399025:
        print "You win!"
    else:
        print "You lose!"
else:
    print "You lose!"
```

Fuzzing Wins

Symbolic Execution Wins

The Goal





American Fuzzy Lop + angr



AFL

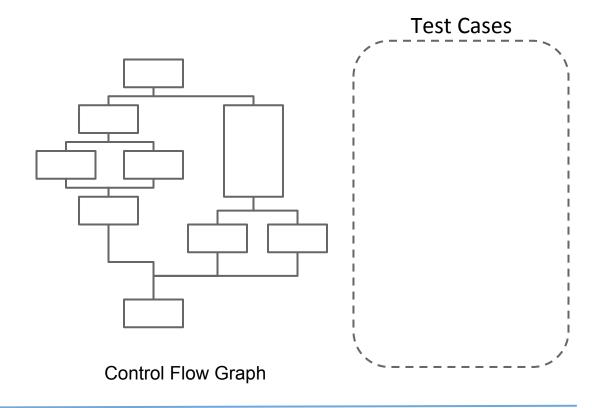
- state-of-the-art instrumented fuzzer
- path uniqueness tracking
- genetic mutations
- open source

angr

- binary analysis platform
- implements symbolic execution engine
- works on binary code
- available on github

Combining the Two (High-level)

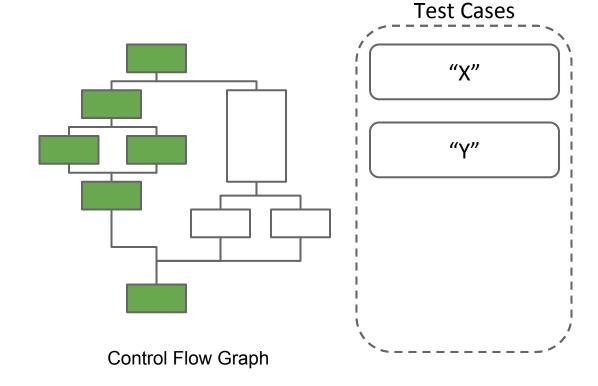




Combining the Two

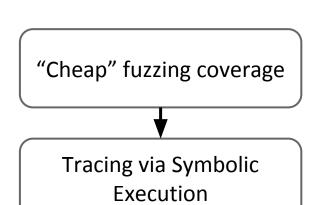


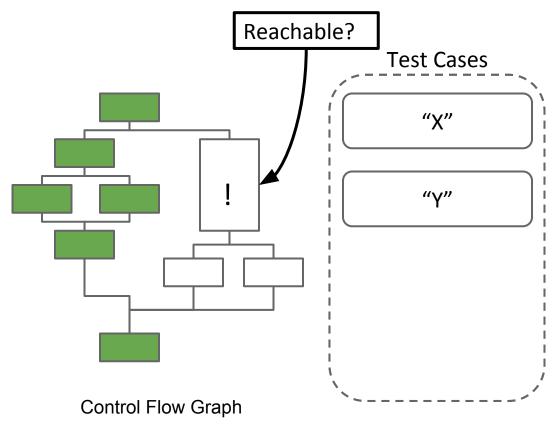
"Cheap" fuzzing coverage



Combining the Two



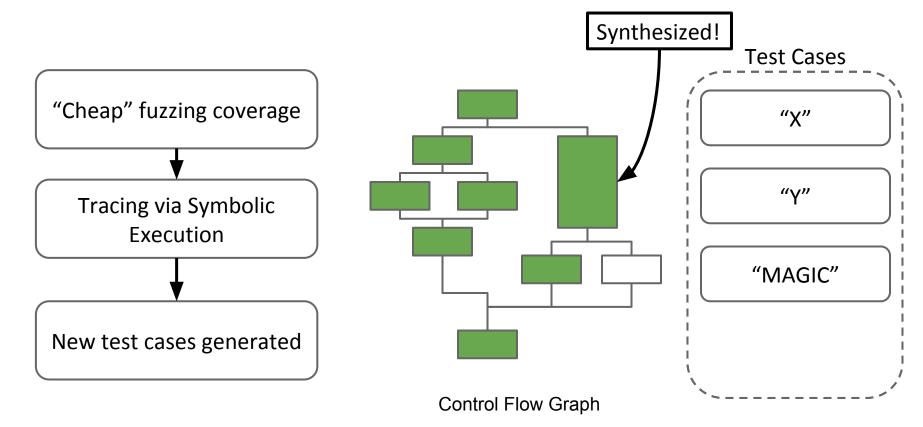




37

Combining the Two





Combining the Two



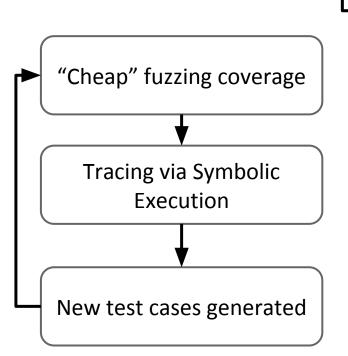
Test Cases

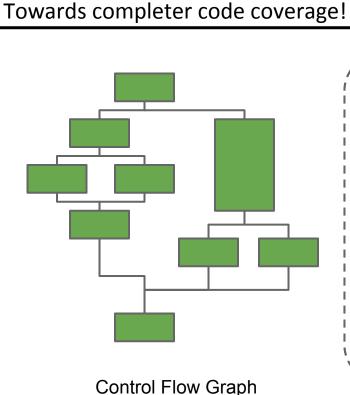
"X"

"V"

"MAGIC"

"MAGICY"



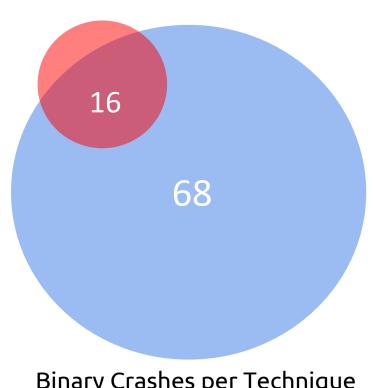


CGC Fuzzer Optimizations



- Symbolic Execution (angr) 16
- Fuzzing (AFL) 68 total
- S & F Shared 13 total

71 / 128 binaries



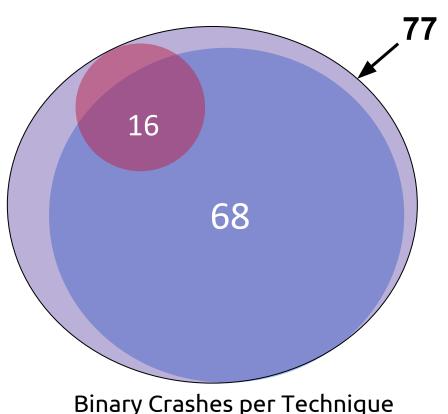
Binary Crashes per Technique

Driller Results



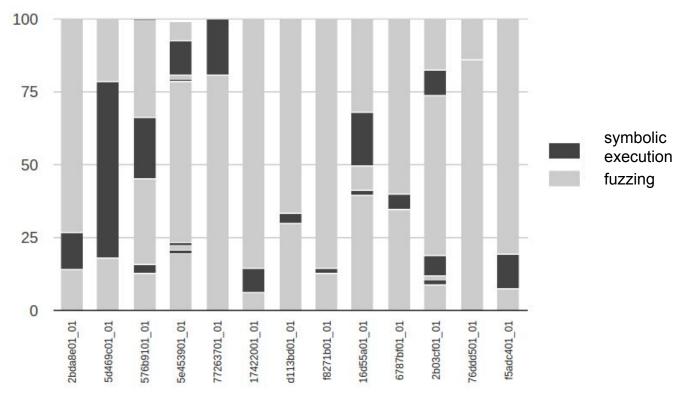
- Symbolic Execution (angr) 16
- Fuzzing (AFL) 68 total
- S & F Shared 13 total
- Driller 77 total

77 / 128 binaries



Binary Crashes per Technique

Distribution of Transitions Found as Iterations of Symbolic Execution and Fuzzing



Roadmap



- DARPA Cyber Grand Challenge (CGC)
 The (almost-)Million Dollar Baby
- Our Cyber Reasoning System (CRS)
 Fancy term for auto-playing a CTF
- Automated Vulnerability Discovery
 Driller
- Automated Vulnerability Exploitation
 How it works
- Auto-exploitation demo using angr
 Open-source binary analysis framework

Previous Work



- Automatic Generation of Control Flow Hijacking Exploits ...
 - Sean Heelan
- AEG: Automatic Exploit Generation
 - Thanassis Avgerinos, CMU
- Unleashing MAYHEM on Binary Code
 - Sang Kil Cha, CMU
- Q: Exploit Hardening Made Easy
 - Edward J. Schwartz, CMU
- More work being done by Julien Vanegue, as far as I know still largely theoretical at this point



- Find a path with safety violations

- Exploit it



```
typedef struct component {
     char name[32];
     int (*do something)(int arg);
} comp t;
comp_t *initialize_component(char *cmp_name) {
     int i = 0:
     struct component *cmp;
     cmp = malloc(sizeof(struct component));
     cmp->do something = sample func;
     while (*cmp_name)
           cmp->name[i++] = *cmp name++;
     cmp->name[i] = '\0';
     return cmp;
 = get input();
cmp = initialize_component(x);
cmp->do_something(1);
```

HEAP

```
Symbolic Byte[0]
Symbolic Byte[1]
Symbolic Byte[2]
Symbolic Byte[3]
Symbolic Byte[4]
Symbolic Byte[5]
Symbolic Byte[6]
Symbolic Byte[7]
...

Symbolic Byte[32] ...
Symbolic Byte[36]
```

call <symbolic byte[36:32]>



Turning the state into an **exploited** state

```
assert state.se.symbolic(state.regs.pc)
```

find_symbolic_buffer is not included in the angr project, but an implementation of this will be shown in the demo

Constrain **buffer** to contain our shellcode

```
buf_addr = find_symbolic_buffer(state, len(shellcode))
mem = state.memory.load(buf_addr, len(shellcode))
state.add_constraints(mem == state.se.bvv(shellcode))
```



Constrain **PC** to point to the buffer

```
state.se.add_constraints(state.regs.pc == buf_addr)
```

Synthesize!

```
exploit = state.posix.dumps(0)
```



Vulnerable Symbolic State (PC hijack)

Constraints to add shellcode to the address space

+

Constraints to make PC point to shellcode

Exploit

Auto Exploitation - Advances



- ROP chain generation and ROP placement (angrop)
 - state-aware ROP
- Write primitives and Read primitives

Forced leaking

Shellcode distribution

Roadmap



- DARPA Cyber Grand Challenge (CGC)
 The (almost-)Million Dollar Baby
- Our Cyber Reasoning System (CRS)
 Fancy term for auto-playing a CTF
- Automated Vulnerability Discovery
 Driller
- Automated Vulnerability Exploitation
 How it works
- Auto-exploitation demo using angr Open-source binary analysis framework

Demo



Available at:

https://github.com/angr/angr-doc/

examples/insomnihack_aeg